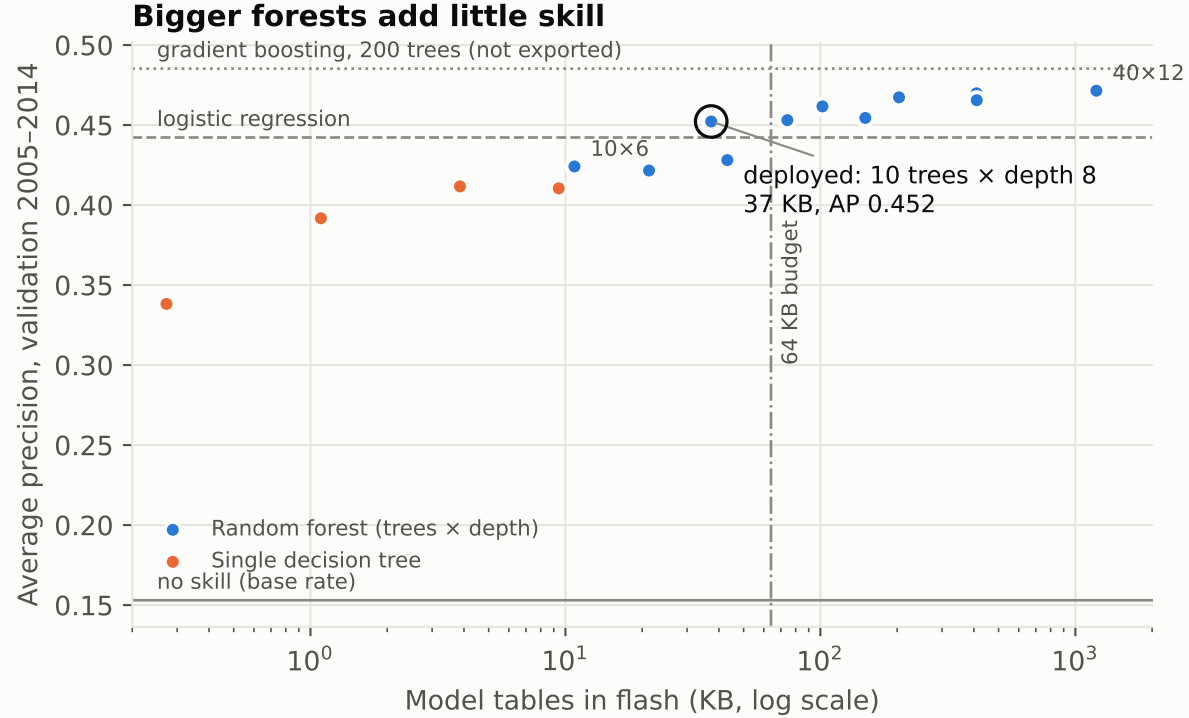
Bigger forests add little skill
gradient boosting, 200 trees (not exported)
logistic regression
40×12
10×6
deployed: 10 trees × depth 8
37 KB, AP 0.452
64 KB budget
Random forest (trees × depth)
Single decision tree
no skill (base rate)
Average precision, validation 2005–2014
Model tables in flash (KB, log scale)
0.50
0.45
0.40
0.35
0.30
0.25
0.20
0.15
10^0
10^1
10^2
10^3